

Classical_ML_Algorithm_IDS_IOT_MEDICAL

May 17, 2025

```
[1]: # importing libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import confusion_matrix, classification_report, \
    ConfusionMatrixDisplay, accuracy_score, recall_score, precision_score, f1_score
from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import AdaBoostClassifier, GradientBoostingClassifier, \
    RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from imblearn.ensemble import BalancedRandomForestClassifier
import xgboost as xgb

[2]: # dataset used
# replace with Zach's df
df2 = pd.read_csv("0.05percent_2classes_processed.csv")
df8 = pd.read_csv("0.05percent_8classes_processed.csv")
df34 = pd.read_csv("0.05percent_34classes_processed.csv")

[3]: # creating dataframes to store result metrics
columns = ["Logistic Regression", "Adaboost", "Gradientboost", "KNN", "XGBoost", \
    "Random Forest", "Random Forest - Bagging"]
index = ["Accuracy", "Recall", "Precision", "F1-Score"]

metrics_2 = pd.DataFrame(index=index, columns=columns)
metrics_8 = pd.DataFrame(index=index, columns=columns)
metrics_34 = pd.DataFrame(index=index, columns=columns)
```

1 Processing and Splitting our dataset Function

80% Train 20% Test

```
[4]: # First, modify your split function to return the scaler
def split(df_name):
    if df_name == "df2":
        label = "benign"
        df = df2.sample(n=400000, random_state=42)
    else:
        label = "label"
        if df_name == "df8": df = df8.sample(n=400000, random_state=42)
        else: df = df34.sample(n=400000, random_state=42)

    # Sorting our dataset into features and target
    X = df.drop(label, axis = 1)
    y = df[label]

    # splitting out dataset to train and test
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳stratify=y)

    # scaling our features
    scaler = StandardScaler()
    scaled_X_train = scaler.fit_transform(X_train)
    scaled_X_test = scaler.transform(X_test) # Note: should use transform, not
↳fit_transform here

    # encoding for XGBoost
    encoder = LabelEncoder()
    encoded_y_train = encoder.fit_transform(y_train)
    encoded_y_test = encoder.transform(y_test) # Should use transform instead
↳of fit_transform

    return scaled_X_train, scaled_X_test, y_train, y_test, encoded_y_train,
↳encoded_y_test, scaler
```

2 Evaluation Functions

```
[5]: # displays the Classification Report and Confusion Matrix
# inserts the metrics of the model into the metrics dataframe
def evaluate(model, y_test, target_names, classes, model_name):
    pred = model.predict(scaled_X_test)

    accuracy = accuracy_score(y_test, pred)
    recall = recall_score(y_test, pred, average='macro')
    precision = precision_score(y_test, pred, average='macro')
    f1 = f1_score(y_test, pred, average='macro')
    insert_metrics(classes, model_name, [accuracy, recall, precision, f1])
```

```

    if model_name == "XGBoost":
        print(classification_report(y_test, pred, target_names = target_names,
→digits = 3))

        if classes == 2 or classes == 8:
            plt.figure(figsize = (4, 2), dpi = 300)
            ConfusionMatrixDisplay(np.round(confusion_matrix(y_test, pred,
→normalize="true"), 2),
                                display_labels = target_names).plot()
            plt.xticks(rotation=90);
        else:
            print(classification_report(y_test, pred, digits = 3))

            if classes == 2 or classes == 8:
                plt.figure(figsize = (4, 2), dpi = 300)
                ConfusionMatrixDisplay(np.round(confusion_matrix(y_test, pred,
→normalize="true"), 2),
                                    display_labels = target_names).plot()
                plt.xticks(rotation=90);

```

```

[6]: # inserts the metrics of the model into the metrics dataframe
def insert_metrics(classes, model_name, metrics):
    if classes == 2:
        metrics_2.loc['Accuracy': 'F1-Score', model_name] = metrics
    elif classes == 8:
        metrics_8.loc['Accuracy': 'F1-Score', model_name] = metrics
    else:
        metrics_34.loc['Accuracy': 'F1-Score', model_name] = metrics

```

```

[7]: # inspired from Zach
def get_weights(y_train):
    import math

    # Get counts in log10
    y_train_count = y_train.value_counts()
    y_log = y_train_count.apply(lambda x: math.log10(x))

    # Find the factor needed to bring minority classes to majority
    y_log_max = y_log.max()
    y_log = y_log.apply(lambda x: math.pow(10, y_log_max - x))

    weight_dict = y_log.to_dict()

    return weight_dict

```

3 Save models

```
[8]: import os
import pickle

def save_model(model, model_name, class_type):
    """
    Save a trained model to disk based on classification type
    """
    # Create directory structure if it doesn't exist
    base_dir = os.path.join("saved_models", f"class_{class_type}")
    os.makedirs(base_dir, exist_ok=True)

    # Create file name with no spaces
    file_name = model_name.replace(" ", "_").lower() + ".pkl"
    file_path = os.path.join(base_dir, file_name)

    # Save the model using pickle
    with open(file_path, 'wb') as f:
        pickle.dump(model, f)

    print(f"Model '{model_name}' saved for {class_type}-class classification at: ↵
    ↵{file_path}")

def save_scaler(scaler, class_type):
    """
    Save the fitted StandardScaler object
    """
    base_dir = os.path.join("saved_models", f"class_{class_type}")
    os.makedirs(base_dir, exist_ok=True)

    file_path = os.path.join(base_dir, "standard_scaler.pkl")

    with open(file_path, 'wb') as f:
        pickle.dump(scaler, f)

    print(f"Scaler saved for {class_type}-class classification at: {file_path}")

def load_model(model_name, class_type):
    """
    Load a saved model from disk
    """
    file_name = model_name.replace(" ", "_").lower() + ".pkl"
    file_path = os.path.join("saved_models", f"class_{class_type}", file_name)

    with open(file_path, 'rb') as f:
        model = pickle.load(f)
```

```

    return model

def load_scaler(class_type):
    """
    Load a saved scaler from disk
    """
    file_path = os.path.join("saved_models", f"class_{class_type}",
↪ "standard_scaler.pkl")

    with open(file_path, 'rb') as f:
        scaler = pickle.load(f)

    return scaler

# Example of how to save models after training
# For 2-class classification models
def save_models_for_class_2():
    # Save the scaler first
    # Make sure to get the scaler that was used to transform the training data
    save_scaler(scaler, 2)

    # Save all models
    save_model(logreg_model, "Logistic Regression", 2)
    save_model(clf, "Adaboost", 2)
    save_model(gbc, "Gradientboost", 2)
    save_model(knn, "KNN", 2)
    save_model(xgc, "XGBoost", 2)
    save_model(rfc, "Random Forest", 2)
    save_model(rfc, "Random Forest Bagging", 2) # Make sure to rename variable

# Similar functions for 8-class and 34-class models
def save_models_for_class_8():
    save_scaler(scaler, 8)
    # Save all the 8-class models using the same pattern as above
    save_model(logreg_model, "Logistic Regression", 8)
    save_model(clf, "Adaboost", 8)
    save_model(gbc, "Gradientboost", 8)
    save_model(knn, "KNN", 8)
    save_model(xgc, "XGBoost", 8)
    save_model(rfc, "Random Forest", 8)
    save_model(rfcb, "Random Forest Bagging", 8)

def save_models_for_class_34():
    save_scaler(scaler, 34)
    # Save all the 34-class models using the same pattern as above

```

```

save_model(logreg_model, "Logistic Regression", 34)
save_model(clf, "Adaboost", 34)
save_model(gbc, "Gradientboost", 34)
save_model(knn, "KNN", 34)
save_model(xgc, "XGBoost", 34)
save_model(rfc, "Random Forest", 34)
save_model(rfcb, "Random Forest Bagging", 34)

# Example of how to load and use a saved model for prediction
def load_and_predict(model_name, class_type, X_new):
    """
    Load a model and scaler, and make predictions on new data

    Parameters:
    -----
    model_name : str
        Name of the model to load
    class_type : int
        Classification type (2, 8, or 34)
    X_new : array-like
        New data to predict on

    Returns:
    -----
    predictions : array
        Predicted classes
    """
    # Load model and scaler
    model = load_model(model_name, class_type)
    scaler = load_scaler(class_type)

    # Scale the new data
    X_new_scaled = scaler.transform(X_new)

    # Make predictions
    predictions = model.predict(X_new_scaled)

    return predictions

```

4 1-Supervised learning

4.1 2 Classes

```

[9]: scaled_X_train, scaled_X_test, y_train, y_test, encoded_y_train, encoded_y_test, \
    ↪ scaler = split("df2")

```

```
[10]: weight_dict = get_weights(y_train)
weight_dict
```

```
[10]: {False: 1.0, True: 41.6552919221541}
```

```
[11]: # essential for XGBoost
target_names = ["Attack", "Benign"]
```

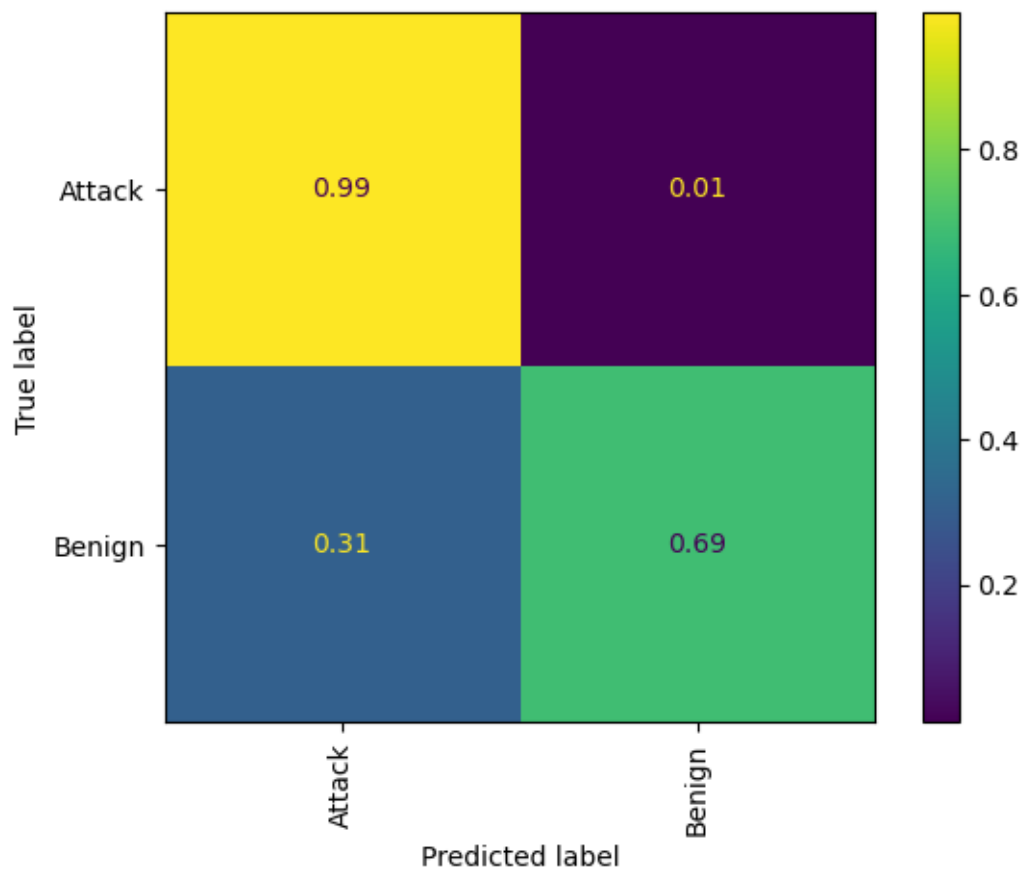
4.1.1 Logistic Regression

```
[12]: logreg_model = LogisticRegression(C = 1, max_iter=500, solver='lbfgs',
    ↪penalty='l2', random_state=42)
logreg_model.fit(scaled_X_train, y_train)

evaluate(logreg_model, y_test, target_names, 2, "Logistic Regression")
```

	precision	recall	f1-score	support
False	0.993	0.994	0.994	78124
True	0.747	0.695	0.720	1876
accuracy			0.987	80000
macro avg	0.870	0.844	0.857	80000
weighted avg	0.987	0.987	0.987	80000

<Figure size 1200x600 with 0 Axes>




4.1.2 Adaboost

```
[13]: clf = AdaBoostClassifier(n_estimators=100, learning_rate = 1, random_state=42)
      clf.fit(scaled_X_train, y_train)

      evaluate(clf, y_test, target_names, 2, "Adaboost")
```

	precision	recall	f1-score	support
False	0.998	0.997	0.997	78124
True	0.877	0.921	0.898	1876
accuracy			0.995	80000
macro avg	0.937	0.959	0.948	80000
weighted avg	0.995	0.995	0.995	80000

<Figure size 1200x600 with 0 Axes>




4.1.3 Gradientboost

```
[14]: gbc = GradientBoostingClassifier(n_estimators = 100, learning_rate = 1,
    ↪max_depth = None)
gbc.fit(scaled_X_train, y_train)

evaluate(gbc, y_test, target_names, 2, "Gradientboost")
```

	precision	recall	f1-score	support
False	0.998	0.998	0.998	78124
True	0.907	0.912	0.910	1876
accuracy			0.996	80000
macro avg	0.953	0.955	0.954	80000
weighted avg	0.996	0.996	0.996	80000

<Figure size 1200x600 with 0 Axes>




4.1.4 KNN

```
[15]: knn = KNeighborsClassifier(n_neighbors = 8)
      knn.fit(scaled_X_train, y_train)

      evaluate(knn, y_test, target_names, 2, "KNN")
```

	precision	recall	f1-score	support
False	0.996	0.994	0.995	78124
True	0.783	0.840	0.810	1876
accuracy			0.991	80000
macro avg	0.890	0.917	0.903	80000
weighted avg	0.991	0.991	0.991	80000

<Figure size 1200x600 with 0 Axes>




4.1.5 XGBoost

```
[16]: xgc = xgb.XGBClassifier(n_estimators=100, learning_rate = 1, booster = "gbtree",
    random_state = 42,
    max_depth = 13, n_jobs=-1)
xgc.fit(scaled_X_train, encoded_y_train)

evaluate(xgc, encoded_y_test, target_names, 2, "XGBoost")
```

	precision	recall	f1-score	support
Attack	0.996	0.996	0.996	78124
Benign	0.839	0.836	0.837	1876
accuracy			0.992	80000
macro avg	0.917	0.916	0.917	80000
weighted avg	0.992	0.992	0.992	80000

<Figure size 1200x600 with 0 Axes>




4.1.6 Random Forest

```
[17]: # Base RF
rfc = RandomForestClassifier(n_estimators=100,n_jobs=-1)
rfc.fit(scaled_X_train, y_train)

evaluate(rfc, y_test, target_names, 2, "Random Forest")
```

	precision	recall	f1-score	support
False	0.999	0.998	0.998	78124
True	0.925	0.949	0.937	1876
accuracy			0.997	80000
macro avg	0.962	0.973	0.968	80000
weighted avg	0.997	0.997	0.997	80000

<Figure size 1200x600 with 0 Axes>




```
[18]: # Bagging-balanced
rfc = BalancedRandomForestClassifier(n_estimators=100, sampling_strategy="all",
    replacement=True)
rfc.fit(scaled_X_train, y_train)

evaluate(rfc, y_test, target_names, 2, "Random Forest - Bagging")
```

	precision	recall	f1-score	support
False	1.000	0.989	0.995	78124
True	0.688	0.999	0.815	1876
accuracy			0.989	80000
macro avg	0.844	0.994	0.905	80000
weighted avg	0.993	0.989	0.990	80000

<Figure size 1200x600 with 0 Axes>



```
[19]: # exporting and displaying the class 2 metrics
metrics_2.to_csv('metrics_2.csv', index=True)
metrics_2
# Add:
```

```
[19]:
```

	Logistic Regression	Adaboost	Gradientboost	KNN	XGBoost	\
Accuracy	0.987325	0.9951	0.99575	0.990788	0.992387	
Recall	0.844459	0.958733	0.954903	0.916986	0.916244	
Precision	0.869905	0.93737	0.952549	0.889666	0.91733	
F1-Score	0.856703	0.947783	0.953723	0.902836	0.916786	

	Random Forest	Random Forest - Bagging
Accuracy	0.996988	0.98935
Recall	0.973486	0.994027
Precision	0.961723	0.843967
F1-Score	0.967526	0.90465

```
[20]: save_models_for_class_2()
```

```

Scaler saved for 2-class classification at:
saved_models/class_2/standard_scaler.pkl
Model 'Logistic Regression' saved for 2-class classification at:
saved_models/class_2/logistic_regression.pkl
Model 'Adaboost' saved for 2-class classification at:
saved_models/class_2/adaboost.pkl
Model 'Gradientboost' saved for 2-class classification at:
saved_models/class_2/gradientboost.pkl
Model 'KNN' saved for 2-class classification at: saved_models/class_2/knn.pkl
Model 'XGBoost' saved for 2-class classification at:
saved_models/class_2/xgboost.pkl
Model 'Random Forest' saved for 2-class classification at:
saved_models/class_2/random_forest.pkl
Model 'Random Forest Bagging' saved for 2-class classification at:
saved_models/class_2/random_forest_bagging.pkl

```

4.2 8 Classes

```
[21]: scaled_X_train, scaled_X_test, y_train, y_test, encoded_y_train, encoded_y_test, \
      ↪ scaler = split("df8")
```

```
[22]: weight_dict = get_weights(y_train)
      weight_dict
```

```
[22]: {'DDoS': 1.0,
      'DoS': 4.217345664488805,
      'Mirai': 12.836500853853368,
      'Benign': 31.061183684350862,
      'Spoofing': 68.8392909896603,
      'Recon': 95.61797291752158,
      'Web': 1429.5766871165652,
      'BruteForce': 2709.5465116279092}
```

```
[23]: y_test.value_counts()
```

```
[23]: label
DDoS          58255
DoS           13813
Mirai          4538
Benign         1876
Spoofing        846
Recon           609
Web             41
BruteForce      22
Name: count, dtype: int64
```

```
[24]: # essential for XGBoost
```


```
target_names = ["Benign", "BruteForce", "DDoS", "Dos", "Mirai", "Recon",  
↳ "Spoofing", "Web"]
```

4.2.1 Logistic Regression

```
[25]: logreg_model = LogisticRegression(C = 1, max_iter=500, solver='lbfgs',  
↳ penalty='l2', random_state=42,n_jobs=-1)  
logreg_model.fit(scaled_X_train, y_train)  
  
evaluate(logreg_model, y_test, target_names, 8, "Logistic Regression")
```

	precision	recall	f1-score	support
Benign	0.695	0.888	0.780	1876
BruteForce	1.000	0.091	0.167	22
DDoS	0.824	0.991	0.900	58255
Dos	0.700	0.111	0.192	13813
Mirai	0.991	0.987	0.989	4538
Recon	0.573	0.340	0.427	609
Spoofing	0.767	0.404	0.529	846
Web	0.000	0.000	0.000	41
accuracy			0.825	80000
macro avg	0.694	0.476	0.498	80000
weighted avg	0.806	0.825	0.772	80000

<Figure size 1200x600 with 0 Axes>



4.2.2 Adaboost

```
[26]: clf = AdaBoostClassifier(n_estimators=100, learning_rate = 1, random_state=42)
      clf.fit(scaled_X_train, y_train)

      evaluate(clf, y_test, target_names, 8, "Adaboost")
```


	precision	recall	f1-score	support
Benign	0.767	0.916	0.835	1876
BruteForce	0.000	0.000	0.000	22
DDoS	0.998	0.997	0.998	58255
DoS	0.995	0.991	0.993	13813
Mirai	0.993	0.981	0.987	4538
Recon	0.521	0.558	0.539	609
Spoofing	0.754	0.597	0.666	846
Web	0.000	0.000	0.000	41

accuracy			0.985	80000
macro avg	0.629	0.630	0.627	80000
weighted avg	0.985	0.985	0.985	80000

/home/faissalm/anaconda3/lib/python3.12/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

<Figure size 1200x600 with 0 Axes>




4.2.3 Gradientboost

```
[27]: gbc = GradientBoostingClassifier(n_estimators=100, learning_rate = 1, max_depth=
      ↳= None)
      gbc.fit(scaled_X_train, y_train)

      evaluate(gbc, y_test, target_names, 8, "Gradientboost")
```

	precision	recall	f1-score	support
Benign	0.936	0.832	0.881	1876
BruteForce	0.778	0.636	0.700	22
DDoS	0.996	1.000	0.998	58255
DoS	0.999	0.999	0.999	13813
Mirai	0.996	0.999	0.998	4538
Recon	0.785	0.772	0.778	609
Spoofing	0.824	0.823	0.823	846
Web	0.405	0.366	0.385	41
accuracy			0.992	80000
macro avg	0.840	0.803	0.820	80000
weighted avg	0.991	0.992	0.992	80000

<Figure size 1200x600 with 0 Axes>



5 KNN


```
[28]: knn = KNeighborsClassifier(n_neighbors = 8)
      knn.fit(scaled_X_train, y_train)

      evaluate(knn, y_test, target_names, 8, "KNN")
```

	precision	recall	f1-score	support
Benign	0.736	0.927	0.820	1876
BruteForce	1.000	0.091	0.167	22
DDoS	0.958	0.985	0.971	58255
DoS	0.928	0.822	0.872	13813
Mirai	0.998	0.990	0.994	4538
Recon	0.692	0.524	0.596	609
Spoofing	0.744	0.491	0.591	846
Web	0.000	0.000	0.000	41

accuracy			0.946	80000
macro avg	0.757	0.604	0.626	80000
weighted avg	0.945	0.946	0.944	80000

<Figure size 1200x600 with 0 Axes>



5.1 XGBoost

```
[29]: # Required to be trained on scaled_X_train, encoded_y_train
# evaluate parameter for y_test to be encoded y_test
xgc = xgb.XGBClassifier(n_estimators=100, learning_rate = 1, booster = "gbtree",
    ↪random_state = 42,
                        max_depth = 13, max_samples = 0.9 ,n_jobs=-1)
xgc.fit(scaled_X_train, encoded_y_train)

evaluate(xgc, encoded_y_test, target_names, 8, "XGBoost")
```

```
/home/faissalm/anaconda3/lib/python3.12/site-packages/xgboost/training.py:183:
UserWarning: [14:28:30] WARNING: /workspace/src/learner.cc:738:
Parameters: { "max_samples" } are not used.
```


```
bst.update(dtrain, iteration=i, fobj=obj)

      precision    recall  f1-score   support

   Benign       0.896     0.933     0.914       1876
  BruteForce    1.000     0.136     0.240         22
     DDoS       1.000     1.000     1.000      58255
       Dos      0.999     0.999     0.999     13813
     Mirai      1.000     0.999     0.999     4538
     Recon      0.796     0.788     0.792         609
   Spoofing      0.822     0.798     0.810         846
       Web      0.357     0.122     0.182          41

 accuracy                   0.994      80000
  macro avg       0.859     0.722     0.742      80000
weighted avg       0.993     0.994     0.993      80000
```

<Figure size 1200x600 with 0 Axes>



5.2 Random Forest


```
[30]: # Base RF
rfc = RandomForestClassifier(n_estimators=100 ,n_jobs=-1)
rfc.fit(scaled_X_train, y_train)

evaluate(rfc, y_test, target_names, 8, "Random Forest")
```

	precision	recall	f1-score	support
Benign	0.895	0.962	0.927	1876
BruteForce	1.000	0.045	0.087	22
DDoS	1.000	1.000	1.000	58255
Dos	1.000	0.999	0.999	13813
Mirai	1.000	1.000	1.000	4538
Recon	0.827	0.801	0.814	609
Spoofing	0.894	0.820	0.856	846
Web	1.000	0.098	0.178	41

accuracy			0.995	80000
macro avg	0.952	0.716	0.733	80000
weighted avg	0.995	0.995	0.994	80000

<Figure size 1200x600 with 0 Axes>




```
[31]: # Bagging-balanced
rfcb = BalancedRandomForestClassifier(n_estimators=100, sampling_strategy="all",
    ↪replacement=True, n_jobs=-1)
rfcb.fit(scaled_X_train, y_train)

evaluate(rfcb, y_test, target_names, 8, "Random Forest - Bagging")
```

precision	recall	f1-score	support
-----------	--------	----------	---------

Benign	0.896	0.792	0.841	1876
BruteForce	0.039	0.727	0.074	22
DDoS	1.000	0.976	0.987	58255
DoS	0.954	0.980	0.967	13813
Mirai	0.997	0.985	0.991	4538
Recon	0.395	0.639	0.488	609
Spoofing	0.509	0.661	0.575	846
Web	0.067	0.561	0.120	41
accuracy			0.966	80000
macro avg	0.607	0.790	0.630	80000
weighted avg	0.979	0.966	0.972	80000

<Figure size 1200x600 with 0 Axes>



```
[32]: # exporting and displaying the class 8 metrics
metrics_8.to_csv('metrics_8.csv', index=True)
```

```
metrics_8
```

```
[32]:
```

	Logistic Regression	Adaboost	Gradientboost	KNN	XGBoost	\
Accuracy	0.82475	0.985012	0.9917	0.946163	0.9936	
Recall	0.476462	0.629999	0.80335	0.603714	0.721853	
Precision	0.693832	0.628506	0.839952	0.756951	0.858733	
F1-Score	0.497933	0.627205	0.820249	0.626488	0.741981	

	Random Forest	Random Forest - Bagging
Accuracy	0.9947	0.966337
Recall	0.71562	0.789926
Precision	0.951963	0.607234
F1-Score	0.732571	0.630485

```
[33]: save_models_for_class_8()
```

```
Scaler saved for 8-class classification at:
saved_models/class_8/standard_scaler.pkl
Model 'Logistic Regression' saved for 8-class classification at:
saved_models/class_8/logistic_regression.pkl
Model 'Adaboost' saved for 8-class classification at:
saved_models/class_8/adaboost.pkl
Model 'Gradientboost' saved for 8-class classification at:
saved_models/class_8/gradientboost.pkl
Model 'KNN' saved for 8-class classification at: saved_models/class_8/knn.pkl
Model 'XGBoost' saved for 8-class classification at:
saved_models/class_8/xgboost.pkl
Model 'Random Forest' saved for 8-class classification at:
saved_models/class_8/random_forest.pkl
Model 'Random Forest Bagging' saved for 8-class classification at:
saved_models/class_8/random_forest_bagging.pkl
```

6 34 Classes

```
[34]: scaled_X_train, scaled_X_test, y_train, y_test, encoded_y_train, encoded_y_test, \
      ↪ scaler = split("df34")
```

```
[35]: weight_dict = get_weights(y_train)
      weight_dict
```

```
[35]: {'DDoS-ICMP_Flood': 1.0,
      'DDoS-UDP_Flood': 1.3327485222273208,
      'DDoS-TCP_Flood': 1.6057560975609746,
      'DDoS-PSHACK_Flood': 1.7570010319182976,
      'DDoS-SYN_Flood': 1.7736628470850218,
      'DDoS-RSTFINFlood': 1.7863680764082344,
      'DDoS-SynonymousIP_Flood': 1.9886825889081294,
```

```

'DoS-UDP_Flood': 2.1790379523389207,
'DoS-TCP_Flood': 2.6990816661200356,
'DoS-SYN_Flood': 3.578302775563443,
'BenignTraffic': 6.580967612954817,
'Mirai-greeth_flood': 7.218859649122804,
'Mirai-udpplain': 7.982056256062071,
'Mirai-greip_flood': 9.630778232884722,
'DDoS-ICMP_Fragmentation': 15.928064516129023,
'MITM-ArpSpoofing': 23.291037735849052,
'DDoS-UDP_Fragmentation': 24.299704724409448,
'DDoS-ACK_Fragmentation': 25.452061855670088,
'DNS_Spoofing': 39.03320158102764,
'Recon-HostDiscovery': 52.58466453674118,
'Recon-OSScan': 73.04289940828397,
'Recon-PortScan': 87.85943060498215,
'DoS-HTTP_Flood': 98.95190380761521,
'VulnerabilityScan': 201.5387755102039,
'DDoS-HTTP_Flood': 259.87894736842094,
'DDoS-SlowLoris': 288.7543859649122,
'DictionaryBruteForce': 574.1511627906976,
'BrowserHijacking': 1097.26666666666653,
'SqlInjection': 1204.317073170731,
'CommandInjection': 1299.3947368421045,
'XSS': 2244.4090909090896,
'Backdoor_Malware': 3291.7999999999997,
'Recon-PingSweep': 3291.7999999999997,
'Uploading_Attack': 24688.499999999998}

```

```
[36]: y_test.value_counts()
```

```

[36]: label
DDoS-ICMP_Flood      12344
DDoS-UDP_Flood       9262
DDoS-TCP_Flood       7687
DDoS-PSHACK_Flood    7026
DDoS-SYN_Flood       6960
DDoS-RSTFINFlood     6910
DDoS-SynonymousIP_Flood 6207
DoS-UDP_Flood        5665
DoS-TCP_Flood        4574
DoS-SYN_Flood        3450
BenignTraffic        1875
Mirai-greeth_flood   1710
Mirai-udpplain       1546
Mirai-greip_flood    1282
DDoS-ICMP_Fragmentation 775
MITM-ArpSpoofing     530

```

DDoS-UDP_Fragmentation	508
DDoS-ACK_Fragmentation	485
DNS_Spoofing	316
Recon-HostDiscovery	235
Recon-OSScan	169
Recon-PortScan	140
DoS-HTTP_Flood	125
VulnerabilityScan	61
DDoS-HTTP_Flood	48
DDoS-SlowLoris	43
DictionaryBruteForce	22
BrowserHijacking	11
SqlInjection	10
CommandInjection	9
XSS	6
Backdoor_Malware	4
Recon-PingSweep	4
Uploading_Attack	1
Name: count, dtype: int64	

```
[37]: # essential for XGBoost
target_names = list(y_test.unique())
```

6.1 Logistic Regression

```
[38]: logreg_model = LogisticRegression(C = 1, max_iter=500, solver='lbfgs',
    ↪penalty='l2', random_state=42, n_jobs=-1)
logreg_model.fit(scaled_X_train, y_train)

evaluate(logreg_model, y_test, target_names, 34, "Logistic Regression")
```

```
/home/faissalm/anaconda3/lib/python3.12/site-
packages/sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/home/faissalm/anaconda3/lib/python3.12/site-
packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/home/faissalm/anaconda3/lib/python3.12/site-
```

packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

	precision	recall	f1-score	support
Backdoor_Malware	0.000	0.000	0.000	4
BenignTraffic	0.674	0.927	0.781	1875
BrowserHijacking	0.000	0.000	0.000	11
CommandInjection	0.000	0.000	0.000	9
DDoS-ACK_Fragmentation	0.884	0.157	0.266	485
DDoS-HTTP_Flood	0.327	0.375	0.350	48
DDoS-ICMP_Flood	0.972	0.996	0.984	12344
DDoS-ICMP_Fragmentation	0.612	0.981	0.754	775
DDoS-PSHACK_Flood	0.999	0.983	0.991	7026
DDoS-RSTFINFlood	0.999	0.999	0.999	6910
DDoS-SYN_Flood	0.650	0.967	0.777	6960
DDoS-SlowLoris	0.222	0.093	0.131	43
DDoS-SynonymousIP_Flood	0.748	0.683	0.714	6207
DDoS-TCP_Flood	0.620	0.945	0.748	7687
DDoS-UDP_Flood	0.722	0.956	0.823	9262
DDoS-UDP_Fragmentation	0.921	0.870	0.895	508
DNS_Spoofing	0.327	0.114	0.169	316
DictionaryBruteForce	1.000	0.045	0.087	22
DoS-HTTP_Flood	0.595	0.600	0.598	125
DoS-SYN_Flood	0.623	0.120	0.201	3450
DoS-TCP_Flood	0.663	0.137	0.227	4574
DoS-UDP_Flood	0.803	0.301	0.437	5665
MITM-ArpSpoofing	0.740	0.472	0.576	530
Mirai-greeth_flood	0.591	0.951	0.729	1710
Mirai-greip_flood	0.536	0.117	0.192	1282
Mirai-udpplain	0.995	0.957	0.976	1546
Recon-HostDiscovery	0.466	0.345	0.396	235
Recon-OSScan	0.128	0.036	0.056	169
Recon-PingSweep	0.000	0.000	0.000	4
Recon-PortScan	0.250	0.029	0.051	140
SqlInjection	0.000	0.000	0.000	10
Uploading_Attack	0.000	0.000	0.000	1
VulnerabilityScan	0.000	0.000	0.000	61
XSS	0.000	0.000	0.000	6
accuracy			0.784	80000
macro avg	0.502	0.416	0.409	80000
weighted avg	0.786	0.784	0.745	80000

/home/faissalm/anaconda3/lib/python3.12/site-

packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning:

Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/home/faissalm/anaconda3/lib/python3.12/site-
packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

6.2 Adaboost

```
[39]: clf = AdaBoostClassifier(n_estimators=100, learning_rate = 1, random_state=42 )
      clf.fit(scaled_X_train, y_train)

      evaluate(clf, y_test, target_names, 34, "Adaboost")
```

```
/home/faissalm/anaconda3/lib/python3.12/site-
packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/home/faissalm/anaconda3/lib/python3.12/site-
packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

	precision	recall	f1-score	support
Backdoor_Malware	0.000	0.000	0.000	4
BenignTraffic	0.546	0.946	0.693	1875
BrowserHijacking	0.000	0.000	0.000	11
CommandInjection	0.000	0.000	0.000	9
DDoS-ACK_Fragmentation	0.857	0.918	0.886	485
DDoS-HTTP_Flood	0.000	0.000	0.000	48
DDoS-ICMP_Flood	1.000	0.996	0.998	12344
DDoS-ICMP_Fragmentation	0.976	0.839	0.902	775
DDoS-PSHACK_Flood	1.000	0.980	0.990	7026
DDoS-RSTFINFlood	0.996	0.998	0.997	6910
DDoS-SYN_Flood	0.663	0.957	0.783	6960
DDoS-SlowLoris	0.000	0.000	0.000	43
DDoS-SynonymousIP_Flood	0.553	0.026	0.050	6207
DDoS-TCP_Flood	0.405	0.967	0.570	7687
DDoS-UDP_Flood	0.709	0.982	0.823	9262
DDoS-UDP_Fragmentation	0.950	0.941	0.946	508
DNS_Spoofing	0.000	0.000	0.000	316
DictionaryBruteForce	0.000	0.000	0.000	22
DoS-HTTP_Flood	0.000	0.000	0.000	125
DoS-SYN_Flood	1.000	0.008	0.016	3450

DoS-TCP_Flood	0.494	0.037	0.070	4574
DoS-UDP_Flood	0.700	0.321	0.440	5665
MITM-ArpSpoofing	0.140	0.045	0.068	530
Mirai-greeth_flood	0.590	0.926	0.721	1710
Mirai-greip_flood	0.575	0.048	0.088	1282
Mirai-udpplain	0.966	0.920	0.942	1546
Recon-HostDiscovery	0.000	0.000	0.000	235
Recon-OSScan	0.000	0.000	0.000	169
Recon-PingSweep	0.000	0.000	0.000	4
Recon-PortScan	0.000	0.000	0.000	140
SqlInjection	0.000	0.000	0.000	10
Uploading_Attack	0.000	0.000	0.000	1
VulnerabilityScan	0.000	0.000	0.000	61
XSS	0.000	0.000	0.000	6
accuracy			0.724	80000
macro avg	0.386	0.349	0.323	80000
weighted avg	0.746	0.724	0.657	80000

```

/home/faissalm/anaconda3/lib/python3.12/site-
packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/home/faissalm/anaconda3/lib/python3.12/site-
packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

```

6.3 Gradientboost

```

[40]: gbc = GradientBoostingClassifier(n_estimators=100, learning_rate = 1, max_depth_
      ↪= None)
      gbc.fit(scaled_X_train, y_train)

      evaluate(gbc, y_test, target_names, 34, "Gradientboost")

```

	precision	recall	f1-score	support
Backdoor_Malware	0.667	0.500	0.571	4
BenignTraffic	0.941	0.862	0.900	1875
BrowserHijacking	0.389	0.636	0.483	11
CommandInjection	0.500	0.444	0.471	9
DDoS-ACK_Fragmentation	0.996	0.990	0.993	485
DDoS-HTTP_Flood	0.935	0.896	0.915	48
DDoS-ICMP_Flood	0.998	1.000	0.999	12344
DDoS-ICMP_Fragmentation	0.997	0.991	0.994	775

DDoS-PSHACK_Flood	0.998	1.000	0.999	7026
DDoS-RSTFINFlood	1.000	1.000	1.000	6910
DDoS-SYN_Flood	0.999	1.000	0.999	6960
DDoS-SlowLoris	1.000	0.953	0.976	43
DDoS-SynonymousIP_Flood	1.000	1.000	1.000	6207
DDoS-TCP_Flood	0.992	1.000	0.996	7687
DDoS-UDP_Flood	1.000	1.000	1.000	9262
DDoS-UDP_Fragmentation	0.994	0.996	0.995	508
DNS_Spoofing	0.683	0.668	0.675	316
DictionaryBruteForce	0.630	0.773	0.694	22
DoS-HTTP_Flood	0.909	0.960	0.934	125
DoS-SYN_Flood	1.000	0.999	1.000	3450
DoS-TCP_Flood	0.999	0.999	0.999	4574
DoS-UDP_Flood	0.975	0.999	0.987	5665
MITM-ArpSpoofing	0.830	0.753	0.789	530
Mirai-greeth_flood	0.999	0.998	0.999	1710
Mirai-greip_flood	0.999	0.998	0.999	1282
Mirai-udpplain	1.000	0.998	0.999	1546
Recon-HostDiscovery	0.787	0.753	0.770	235
Recon-OSScan	0.635	0.586	0.609	169
Recon-PingSweep	0.000	0.000	0.000	4
Recon-PortScan	0.594	0.607	0.601	140
SqlInjection	0.231	0.300	0.261	10
Uploading_Attack	1.000	1.000	1.000	1
VulnerabilityScan	1.000	0.902	0.948	61
XSS	0.500	0.167	0.250	6
accuracy				0.990 80000
macro avg				0.829 0.816 0.818 80000
weighted avg				0.990 0.990 0.990 80000

6.4 KNN

```
[41]: knn = KNeighborsClassifier(n_neighbors = 8, n_jobs=-1)
      knn.fit(scaled_X_train, y_train)

      evaluate(knn, y_test, target_names, 34, "KNN")
```

```
/home/faissalm/anaconda3/lib/python3.12/site-
packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/home/faissalm/anaconda3/lib/python3.12/site-
packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
```



```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

	precision	recall	f1-score	support
Backdoor_Malware	0.000	0.000	0.000	4
BenignTraffic	0.714	0.958	0.818	1875
BrowserHijacking	0.000	0.000	0.000	11
CommandInjection	0.000	0.000	0.000	9
DDoS-ACK_Fragmentation	0.690	0.625	0.656	485
DDoS-HTTP_Flood	0.542	0.667	0.598	48
DDoS-ICMP_Flood	0.988	0.993	0.991	12344
DDoS-ICMP_Fragmentation	0.723	0.836	0.776	775
DDoS-PSHACK_Flood	0.999	0.996	0.998	7026
DDoS-RSTFINFlood	0.999	0.999	0.999	6910
DDoS-SYN_Flood	0.901	0.964	0.931	6960
DDoS-SlowLoris	0.686	0.558	0.615	43
DDoS-SynonymousIP_Flood	0.957	0.944	0.951	6207
DDoS-TCP_Flood	0.930	0.970	0.950	7687
DDoS-UDP_Flood	0.899	0.953	0.925	9262
DDoS-UDP_Fragmentation	0.962	0.858	0.907	508
DNS_Spoofing	0.475	0.301	0.368	316
DictionaryBruteForce	0.500	0.045	0.083	22
DoS-HTTP_Flood	0.845	0.696	0.763	125
DoS-SYN_Flood	0.907	0.802	0.851	3450
DoS-TCP_Flood	0.952	0.876	0.912	4574
DoS-UDP_Flood	0.902	0.812	0.855	5665
MITM-ArpSpoofing	0.742	0.417	0.534	530
Mirai-greeth_flood	0.937	0.950	0.943	1710
Mirai-greip_flood	0.921	0.906	0.914	1282
Mirai-udpplain	0.997	0.984	0.991	1546
Recon-HostDiscovery	0.596	0.566	0.581	235
Recon-OSScan	0.375	0.124	0.187	169
Recon-PingSweep	0.000	0.000	0.000	4
Recon-PortScan	0.457	0.150	0.226	140
SqlInjection	0.000	0.000	0.000	10
Uploading_Attack	0.000	0.000	0.000	1
VulnerabilityScan	0.667	0.787	0.722	61
XSS	0.000	0.000	0.000	6
accuracy			0.932	80000
macro avg	0.625	0.581	0.590	80000
weighted avg	0.931	0.932	0.930	80000

```
/home/faissalm/anaconda3/lib/python3.12/site-  
packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning:  
Precision is ill-defined and being set to 0.0 in labels with no predicted  
samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

```
/home/faissalm/anaconda3/lib/python3.12/site-
packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

6.5 XGBoost

```
[42]: xgc = xgb.XGBClassifier(n_estimators=100, learning_rate = 1, booster = "gbtree",
    ↪ random_state = 42,
        max_depth = 13, max_samples = 0.9 , n_jobs=-1)
xgc.fit(scaled_X_train, encoded_y_train)

evaluate(xgc, encoded_y_test, target_names, 34, "XGBoost")
```

```
/home/faissalm/anaconda3/lib/python3.12/site-packages/xgboost/training.py:183:
UserWarning: [14:42:50] WARNING: /workspace/src/learner.cc:738:
Parameters: { "max_samples" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
```

	precision	recall	f1-score	support
BenignTraffic	0.000	0.000	0.000	4
DDoS-UDP_Flood	0.602	0.418	0.494	1875
DoS-UDP_Flood	0.000	0.000	0.000	11
DDoS-PSHACK_Flood	0.000	0.000	0.000	9
DDoS-SYN_Flood	0.000	0.000	0.000	485
DDoS-SynonymousIP_Flood	0.000	0.000	0.000	48
Recon-PortScan	0.961	0.510	0.666	12344
DDoS-ICMP_Flood	0.667	0.005	0.010	775
DDoS-TCP_Flood	0.931	0.849	0.888	7026
DDoS-RSTFINFlood	0.000	0.000	0.000	6910
DNS_Spoofing	0.122	0.217	0.156	6960
Mirai-greip_flood	0.000	0.000	0.000	43
DDoS-ICMP_Fragmentation	0.000	0.000	0.000	6207
DoS-SYN_Flood	0.139	0.876	0.241	7687
DoS-TCP_Flood	0.000	0.000	0.000	9262
Mirai-udpplain	0.000	0.000	0.000	508
DDoS-ACK_Fragmentation	0.000	0.000	0.000	316
Mirai-greeth_flood	0.000	0.000	0.000	22
DoS-HTTP_Flood	0.000	0.000	0.000	125
DDoS-UDP_Fragmentation	0.000	0.000	0.000	3450
MITM-ArpSpoofing	0.054	0.002	0.005	4574
Recon-HostDiscovery	0.072	0.054	0.062	5665
Recon-OSScan	0.000	0.000	0.000	530
SqlInjection	0.000	0.000	0.000	1710
DDoS-HTTP_Flood	0.000	0.000	0.000	1282
DDoS-SlowLoris	0.031	0.011	0.016	1546

CommandInjection	0.263	0.021	0.039	235
VulnerabilityScan	0.000	0.000	0.000	169
DictionaryBruteForce	0.000	0.000	0.000	4
BrowserHijacking	0.000	0.000	0.000	140
XSS	0.000	0.000	0.000	10
Backdoor_Malware	0.000	0.000	0.000	1
Recon-PingSweep	0.000	0.000	0.000	61
Uploading_Attack	0.000	0.000	0.000	6
accuracy			0.270	80000
macro avg	0.113	0.087	0.076	80000
weighted avg	0.284	0.270	0.234	80000

```

/home/faissalm/anaconda3/lib/python3.12/site-
packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/home/faissalm/anaconda3/lib/python3.12/site-
packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/home/faissalm/anaconda3/lib/python3.12/site-
packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/home/faissalm/anaconda3/lib/python3.12/site-
packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

```

6.6 Random Forest

```

[43]: # Base RF
rfc = RandomForestClassifier(n_estimators=100 , n_jobs=-1)
rfc.fit(scaled_X_train, y_train)

evaluate(rfc, y_test, target_names, 34, "Random Forest")

```

```

/home/faissalm/anaconda3/lib/python3.12/site-
packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/home/faissalm/anaconda3/lib/python3.12/site-

```

packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

	precision	recall	f1-score	support
Backdoor_Malware	0.000	0.000	0.000	4
BenignTraffic	0.858	0.989	0.919	1875
BrowserHijacking	0.000	0.000	0.000	11
CommandInjection	0.000	0.000	0.000	9
DDoS-ACK_Fragmentation	0.996	0.986	0.991	485
DDoS-HTTP_Flood	1.000	0.958	0.979	48
DDoS-ICMP_Flood	1.000	1.000	1.000	12344
DDoS-ICMP_Fragmentation	0.995	0.994	0.994	775
DDoS-PSHACK_Flood	1.000	0.999	1.000	7026
DDoS-RSTFINFlood	1.000	1.000	1.000	6910
DDoS-SYN_Flood	0.999	0.999	0.999	6960
DDoS-SlowLoris	0.741	0.930	0.825	43
DDoS-SynonymousIP_Flood	1.000	1.000	1.000	6207
DDoS-TCP_Flood	1.000	0.999	1.000	7687
DDoS-UDP_Flood	0.999	1.000	0.999	9262
DDoS-UDP_Fragmentation	0.998	0.984	0.991	508
DNS_Spoofing	0.765	0.671	0.715	316
DictionaryBruteForce	1.000	0.091	0.167	22
DoS-HTTP_Flood	0.976	0.976	0.976	125
DoS-SYN_Flood	0.998	0.999	0.998	3450
DoS-TCP_Flood	1.000	0.998	0.999	4574
DoS-UDP_Flood	0.999	0.999	0.999	5665
MITM-ArpSpoofing	0.885	0.811	0.846	530
Mirai-greeth_flood	0.998	0.998	0.998	1710
Mirai-greip_flood	0.996	0.995	0.996	1282
Mirai-udpplain	0.999	0.999	0.999	1546
Recon-HostDiscovery	0.827	0.774	0.800	235
Recon-OSScan	0.874	0.533	0.662	169
Recon-PingSweep	0.000	0.000	0.000	4
Recon-PortScan	0.811	0.550	0.655	140
SqlInjection	0.000	0.000	0.000	10
Uploading_Attack	0.000	0.000	0.000	1
VulnerabilityScan	0.951	0.951	0.951	61
XSS	0.000	0.000	0.000	6
accuracy			0.993	80000
macro avg	0.755	0.711	0.719	80000
weighted avg	0.992	0.993	0.992	80000

/home/faissalm/anaconda3/lib/python3.12/site-
packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning:

Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/home/faissalm/anaconda3/lib/python3.12/site-
packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

```
[44]: # Bagging-balanced
rfcb = BalancedRandomForestClassifier(n_estimators=100, sampling_strategy="all",
↪replacement=True)
rfcb.fit(scaled_X_train, y_train)
evaluate(rfcb, y_test, target_names, 34, "Random Forest - Bagging")
```

	precision	recall	f1-score	support
Backdoor_Malware	0.025	0.750	0.049	4
BenignTraffic	0.847	0.309	0.453	1875
BrowserHijacking	0.012	0.273	0.023	11
CommandInjection	0.045	0.444	0.082	9
DDoS-ACK_Fragmentation	0.807	0.930	0.864	485
DDoS-HTTP_Flood	0.221	0.708	0.337	48
DDoS-ICMP_Flood	0.974	0.993	0.984	12344
DDoS-ICMP_Fragmentation	0.997	0.899	0.946	775
DDoS-PSHACK_Flood	0.992	0.996	0.994	7026
DDoS-RSTFINFlood	0.999	0.998	0.998	6910
DDoS-SYN_Flood	0.953	0.949	0.951	6960
DDoS-SlowLoris	0.247	0.558	0.343	43
DDoS-SynonymousIP_Flood	0.975	0.948	0.961	6207
DDoS-TCP_Flood	0.985	0.985	0.985	7687
DDoS-UDP_Flood	0.775	0.996	0.872	9262
DDoS-UDP_Fragmentation	0.744	0.967	0.841	508
DNS_Spoofing	0.667	0.063	0.116	316
DictionaryBruteForce	0.000	0.000	0.000	22
DoS-HTTP_Flood	0.510	0.600	0.551	125
DoS-SYN_Flood	0.921	0.971	0.945	3450
DoS-TCP_Flood	0.903	0.991	0.945	4574
DoS-UDP_Flood	0.975	0.402	0.570	5665
MITM-ArpSpoofing	0.357	0.481	0.410	530
Mirai-greeth_flood	0.976	0.798	0.878	1710
Mirai-greip_flood	0.791	0.950	0.864	1282
Mirai-udpplain	0.968	0.993	0.980	1546
Recon-HostDiscovery	0.610	0.366	0.457	235
Recon-OSScan	0.132	0.101	0.114	169
Recon-PingSweep	0.006	0.250	0.012	4

Recon-PortScan	0.123	0.064	0.085	140
SqlInjection	0.019	0.200	0.034	10
Uploading_Attack	0.006	1.000	0.012	1
VulnerabilityScan	0.229	0.787	0.354	61
XSS	0.023	0.333	0.043	6
accuracy			0.907	80000
macro avg	0.553	0.649	0.531	80000
weighted avg	0.927	0.907	0.902	80000

```
[45]: # exporting and displaying the class 34 metrics
metrics_34.to_csv('metrics_34.csv', index=True)
metrics_34
```

```
[45]:
```

	Logistic Regression	Adaboost	Gradientboost	KNN	XGBoost \
Accuracy	0.783575	0.723513	0.99035	0.932013	0.270375
Recall	0.416319	0.348667	0.815507	0.580547	0.08718
Precision	0.501909	0.385901	0.828673	0.625427	0.113056
F1-Score	0.409058	0.323045	0.817736	0.589563	0.075797

	Random Forest	Random Forest - Bagging
Accuracy	0.99285	0.906525
Recall	0.71123	0.648675
Precision	0.754803	0.553339
F1-Score	0.719293	0.530971

```
[46]: save_models_for_class_34()
```

```
Scaler saved for 34-class classification at:
saved_models/class_34/standard_scaler.pkl
Model 'Logistic Regression' saved for 34-class classification at:
saved_models/class_34/logistic_regression.pkl
Model 'Adaboost' saved for 34-class classification at:
saved_models/class_34/adaboost.pkl
Model 'Gradientboost' saved for 34-class classification at:
saved_models/class_34/gradientboost.pkl
Model 'KNN' saved for 34-class classification at: saved_models/class_34/knn.pkl
Model 'XGBoost' saved for 34-class classification at:
saved_models/class_34/xgboost.pkl
Model 'Random Forest' saved for 34-class classification at:
saved_models/class_34/random_forest.pkl
Model 'Random Forest Bagging' saved for 34-class classification at:
saved_models/class_34/random_forest_bagging.pkl
```

7 Metric Plot

```
[47]: # List of model names
model_names = ["Logistic Regression", "Adaboost", "Gradientboost", "KNN",
               ↪ "XGBoost", "Random Forest", "Random Forest - Bagging"]

# List of dataframes
dataframes = [metrics_2, metrics_8, metrics_34]

# List of classes
classes = [2, 8, 34]

# List of metric names
metrics = ["Accuracy", "Recall", "Precision", "F1-Score"]

# Create a color mapping dictionary for each model
model_colors = {
    "Logistic Regression": 'tab:blue',
    "Adaboost": 'tab:orange',
    "Gradientboost": 'tab:green',
    "KNN": 'tab:red',
    "XGBoost": 'tab:pink',
    "Random Forest": 'tab:gray',
    "Random Forest - Bagging": 'tab:cyan'
}

# Creating subplots for each metric
for metric in metrics:
    fig, axs = plt.subplots(1, len(classes), figsize=(15, 5), sharey=True)
    fig.suptitle(f"{metric} Scores")

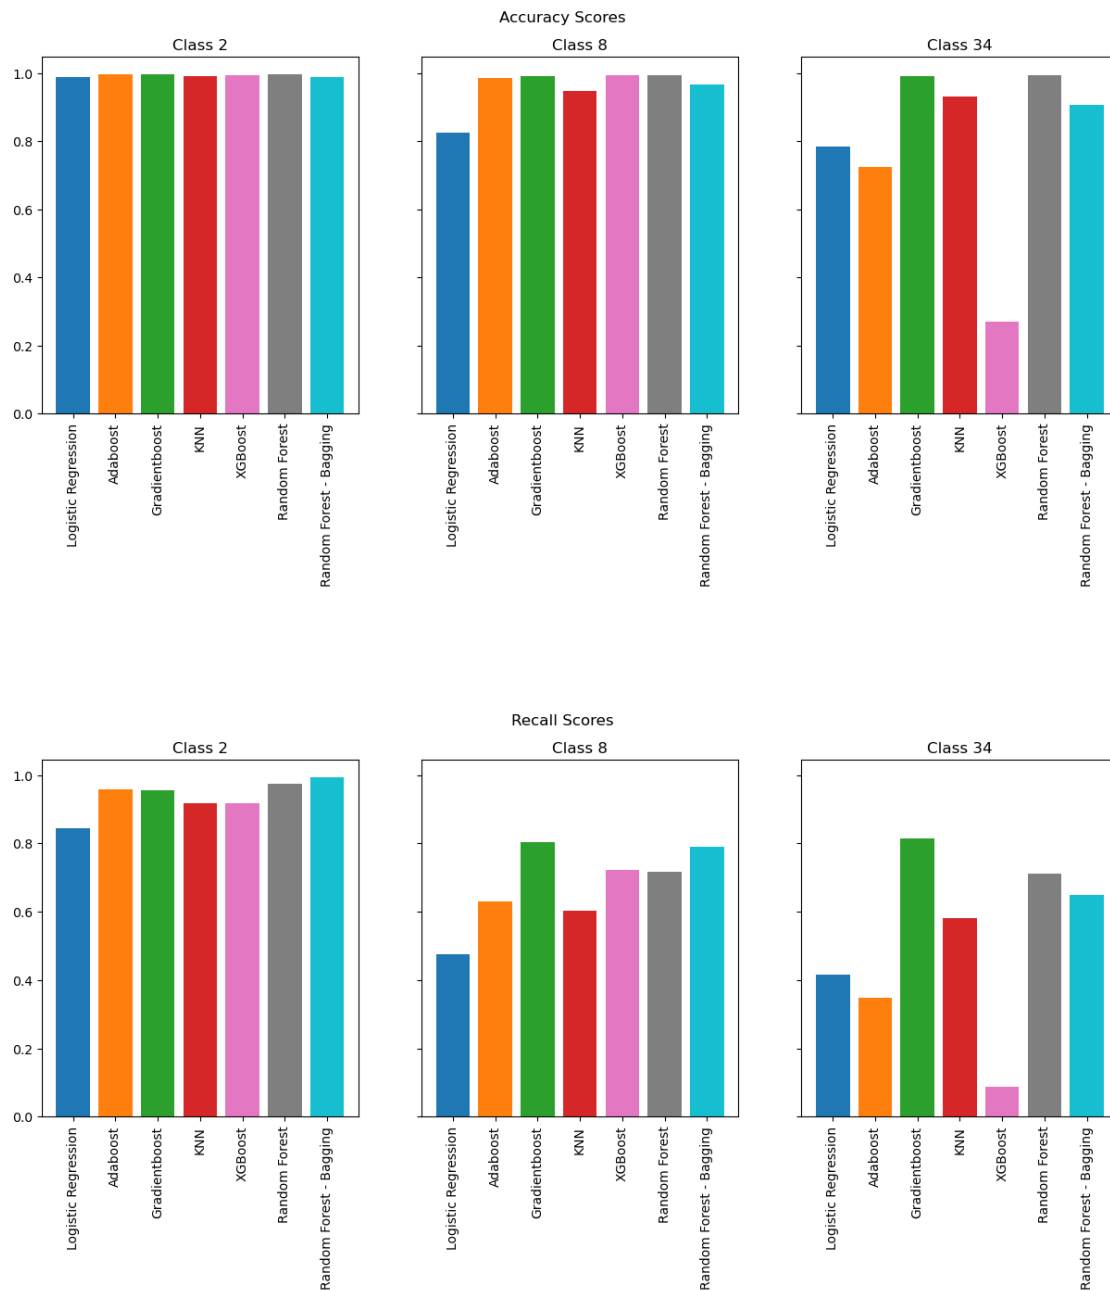
    for i, df in enumerate(dataframes):
        axs[i].set_title(f"Class {classes[i]}")


        # Plotting the respective metric score for each model with specified ↪
        ↪ color
        for model_name in model_names:
            color = model_colors[model_name]
            axs[i].bar(model_name, df.loc[metric, model_name], color=color)

        # Set x-axis ticks and labels, rotating labels by 90 degrees
        axs[i].set_xticks(range(len(model_names)))
        axs[i].set_xticklabels(model_names, rotation=90)

    # Save the figure
    plt.savefig(f'{metric}_scores.png', bbox_inches='tight')
```

```
plt.show()
```





8 Metric Table

```
[48]: # Concatenate them along rows
concatenated_data = pd.concat([metrics_2, metrics_8, metrics_34],
                                ↪ ignore_index=False)

# List of model names
```

```

model_names = ["Logistic Regression", "Adaboost", "Gradientboost", "KNN",
               ↪ "XGBoost", "Random Forest", "Random Forest - Bagging"]

# List of metric names
metrics = ["Accuracy", "Recall", "Precision", "F1-Score"]

# Create a DataFrame with the concatenated data
result_df = pd.DataFrame(concatenated_data.values,
                          pd.MultiIndex.from_product([["2 classes", "8 classes",
               ↪ "34 classes"], metrics]),
                          columns=model_names)

# Display the result
result_df

```

```

[48]:

```

		Logistic Regression	Adaboost	Gradientboost	KNN \
2 classes	Accuracy	0.987325	0.9951	0.99575	0.990788
	Recall	0.844459	0.958733	0.954903	0.916986
	Precision	0.869905	0.93737	0.952549	0.889666
	F1-Score	0.856703	0.947783	0.953723	0.902836
8 classes	Accuracy	0.82475	0.985012	0.9917	0.946163
	Recall	0.476462	0.629999	0.80335	0.603714
	Precision	0.693832	0.628506	0.839952	0.756951
	F1-Score	0.497933	0.627205	0.820249	0.626488
34 classes	Accuracy	0.783575	0.723513	0.99035	0.932013
	Recall	0.416319	0.348667	0.815507	0.580547
	Precision	0.501909	0.385901	0.828673	0.625427
	F1-Score	0.409058	0.323045	0.817736	0.589563

		XGBoost	Random Forest	Random Forest - Bagging
2 classes	Accuracy	0.992387	0.996988	0.98935
	Recall	0.916244	0.973486	0.994027
	Precision	0.91733	0.961723	0.843967
	F1-Score	0.916786	0.967526	0.90465
8 classes	Accuracy	0.9936	0.9947	0.966337
	Recall	0.721853	0.71562	0.789926
	Precision	0.858733	0.951963	0.607234
	F1-Score	0.741981	0.732571	0.630485
34 classes	Accuracy	0.270375	0.99285	0.906525
	Recall	0.08718	0.71123	0.648675
	Precision	0.113056	0.754803	0.553339
	F1-Score	0.075797	0.719293	0.530971